

## High Performance Computing Final Project Proposal

### Rationale:

I played trumpet and wrote a lot of music in high school, so I have always been interested in audio processing. I am also currently enrolled in Digital Signal Processing, where we have covered discrete time convolutions and digital FIR filters in depth using MATLAB. I would like to connect these by completing a project evaluating an optimized audio processing workload in C++. My goal is to discover which optimization techniques (likely a combination of AVX and multi-threading) can produce the fastest execution time for a high-order low-pass FIR filter.

In addition to audio processing, I would also like to accelerate Sobel Filter edge detection using a GPU. I have not worked with GPUs or image processing before, but I think this would be a valuable learning experience to become familiar with CUDA programming, GPU architecture, and edge detection algorithms. If I do not have enough time to implement this workload on a GPU, I will use a CPU implementation with parallelism.

### Workloads:

I will evaluate audio signal processing through convolution using a high-order low-pass FIR filter as well as Sobel Filter edge detection.

### Input Datasets:

For the audio signal processing workload, I plan to use audio samples from the Star Wars movies as datasets since most of the soundtrack is fairly well known, the songs have a wide range of frequencies, and they are long enough to obtain noticeable speedups in parallel. I found MP3 files available here: <https://archive.org/details/13BinarySunsetAlternate> which I can convert into stereo (2-channel) WAV files for processing. For the Sobel Filter edge detection workload, I plan to use PPM files since they are easy to parse in C++. From <https://filesamples.com/formats/ppm> I found an image with noticeable edges available in different sizes ranging from 640x426 to 5184x3456 pixels, which will be good to evaluate weak scaling. This image is shown below:



#### Platforms:

I will evaluate the audio processing workload on CPU nodes available on the Explorer cluster, and possibly the Vector COE server. I will evaluate the Sobel Filter edge detection workload on one or more GPUs available on the Explorer cluster. I will write my programs using C++.

#### Experiments:

For the audio processing workload, I will verify my solutions by outputting and listening to a filtered WAV file to ensure high frequencies have been removed. To design the initial filter and generate FIR coefficients, I will use MATLAB. I can also check my results by testing the same input data in an equivalent MATLAB process. Additionally, I will time my results for different numbers of threads, and with/without AVX support. For the Sobel Filter edge detection algorithm, I will visually compare my results to the original image to ensure edge detection was successful. Finally, I will time the GPU acceleration to evaluate the speedup obtained. I will use the high precision clock available with chrono to time results.

#### Results and Grade Expectations:

A =

- Audio processing workload evaluated with at least 3 WAV input/outputs on the Explorer cluster, including a comparison a serial implementation as well with OpenMP, AVX support, and a combination of the two
- Sobel Filter edge detection workload evaluated with at least 3 different image sizes on one GPU node on the Explorer cluster
- All results reported and analyzed thoroughly in the project writeup, including a strong scalability report for the audio processing workload and a weak scaling scalability report for the Sobel Filter edge detection workload

A- =

- Audio processing workload evaluated with at least 3 array input/outputs (WAV files may be parsed and written outside my C++ program using Python or MATLAB) on the Explorer cluster, including a comparison a serial implementation as well with OpenMP, AVX support, and a combination of the two
- Sobel Filter edge detection workload evaluated with at least 3 different image sizes on one GPU node or one CPU node on the both Explorer and COE Vector, including a comparison of a serial implementation vs a multithread.
- All results reported and analyzed thoroughly in the project writeup, including a strong scalability report for the audio processing workload and a weak scaling scalability report for the Sobel Filter edge detection workload

B+ =

- Audio processing workload evaluated with at least 3 WAV input/outputs on both the Explorer cluster, including a comparison a serial implementation as well with OpenMP, AVX support, and a combination of the two OR Sobel Filter edge detection workload evaluated with at least 3 different image sizes on one GPU node on the Explorer cluster
- All results reported and analyzed thoroughly in the project writeup, including either a strong scalability report for the audio processing workload or a weak scaling scalability report for the Sobel Filter edge detection workload

B =

- Audio processing workload evaluated with at least 3 array input/outputs (WAV files may be parsed and written outside my C++ program using Python or MATLAB) on the Explorer cluster, including a comparison a serial implementation as well with OpenMP, AVX support, and a combination of the two OR Sobel Filter edge detection workload evaluated with at least 3 different image sizes on one CPU node on the both Explorer and COE Vector, including a comparison a serial implementation vs a multithread
- All results reported and analyzed thoroughly in the project writeup, including either a strong scalability report for the audio processing workload or a weak scaling scalability report for the Sobel Filter edge detection workload

B- =

- Audio processing workload evaluated with at least 3 array input/outputs (WAV files may be parsed and written outside my C++ program using Python or MATLAB) on the Explorer cluster, including a comparison a serial implementation vs a multithread implementation OR Sobel Filter edge detection workload evaluated with at least 3 different image sizes on the both Explorer and COE Vector
- All results reported and analyzed thoroughly in the project writeup, including either a strong scalability report for the audio processing workload or a weak scaling scalability report for the Sobel Filter edge detection workload

C+ =

- Audio processing workload evaluated with at least 3 array input/outputs (WAV files may be parsed and written outside my C++ program using Python or MATLAB) on the Explorer cluster, including a comparison a serial implementation vs a multithread implementation OR Sobel Filter edge detection workload evaluated with at least 3 different image sizes on one CPU node on the both Explorer and COE Vector, including a comparison a serial implementation vs a multithread
- All results reported in the project writeup and some analysis included

C =

- Audio processing workload evaluated with at least 3 array input/outputs (WAV files may be parsed and written outside my C++ program using Python or MATLAB) on the Explorer cluster using a multithreaded implementation OR Sobel Filter edge detection workload evaluated with at least 3 different image sizes on one CPU node on the both Explorer and COE Vector
- All results reported in the project writeup and some analysis included

C- =

- Audio processing workload evaluated with one array input/outputs (WAV files may be parsed and written outside my C++ program using Python or MATLAB) on the Explorer cluster using a multithreaded implementation OR Sobel Filter edge detection workload evaluated with at one image size on one CPU node on the both Explorer and COE Vector
- All results reported in the project writeup and some analysis included

Below C- =

- Exploration of either workload, without a fully successful implementation
- Writeup includes details about progress made, challenges faced, and strategies to achieve better results for future work

Extra Credit Opportunities:

- Audio processing workload evaluated on the COE Vector system
- Audio processing workload evaluated using Pthreads
- Multiple FIR filters tested for the audio processing workload
- Sobel Filter edge detection evaluated on one or more CPU nodes to compare to GPU performance
- An additional image processing workload evaluated